

# MACHINE LEARNING IN PHYSICS FOUNDATIONS 2

Harrison B. Prosper  
PHY 6938 Fall 2024

---

# Recap: AI, ML, and DL

## **Artificial Intelligence**

Algorithms that cause machines to exhibit human- or *super-human*-level intelligence.

## **Machine Learning**

Algorithms for modeling data


## **Deep Learning**

ML using (large) neural networks

# Recap: Deep Learning

Uses functions formed by *composition*:

$$\begin{aligned} f(x) &= f_m \circ f_{m-1} \circ \cdots f_1 \\ &= f_m(f_{m-1}(\cdots f_1(x)) \cdots) \end{aligned}$$



**Example:**

$$f(x) = \text{softmax} \left( \text{dropout}(\text{linear}(\text{flatten}(\mathbf{g}(\mathbf{c}(\mathbf{h}(\mathbf{c}(x))))))) \right)$$

# Recap: The Perceptron


$$y = g(xA^T + b)$$

$x$  is matrix (one or more rows) of input data

$A$  is a matrix of **weights**

$b$  is a matrix of **biases**

$g(*)$  is a nonlinear function



**Activations:**  $g(z) = \text{relu}(z) = \max(0, z)$

$g(z) = \tanh(z)$

$g(z) = \text{sigmoid}(z) = 1 / (1 + \exp(-z))$


<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon>

# MULTI-NODE PERCEPTRON

# Multi-Node Perceptron

$$g(z) = \text{relu}(z)$$


$$y_1 = g(xA_1^T + b_1)$$



**Example:**

$$A_1 = (2, -3, 1), \quad b_1 = -5$$


$$y_2 = g(xA_2^T + b_2)$$



$$A_2 = (1, 2, 3), \quad b_2 = -4$$

# Multi-Node Perceptron

A multi-node perceptron is often drawn as follows,



computed sequentially:

$$z = xA^T + b$$

$y = g(z)$  and, usually, the

**activation function**  $g(z)$  is applied *elementwise*, that is, to every element of its matrix input. (Demo)

# Multi-Node Perceptron


Consider the task of modeling the data triplets  $D = \{(x_1, x_2, y)_{i=1}^N\}$  with a function of the form  $y = f(x_1, x_2, \omega)$ . Let's try a 2-input 4-node perceptron with  $g(z) = z\sigma(z)$ ,

where,

$$\sigma(z) = \frac{1}{(1 + e^{-z})}$$

$z\sigma(z)$  is called a **silu**.


$\omega$  are free parameters.






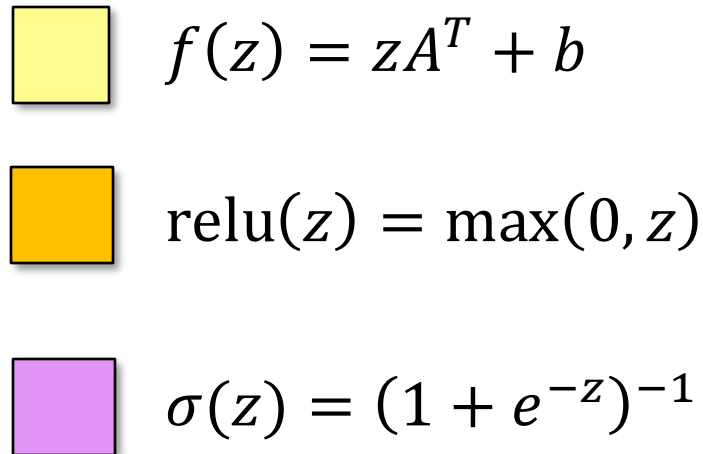
# Multi-Node Perceptron

Then, let's feed the output of the first perceptron into a 4-input 1-node perceptron:



# Multi-Node Perceptron

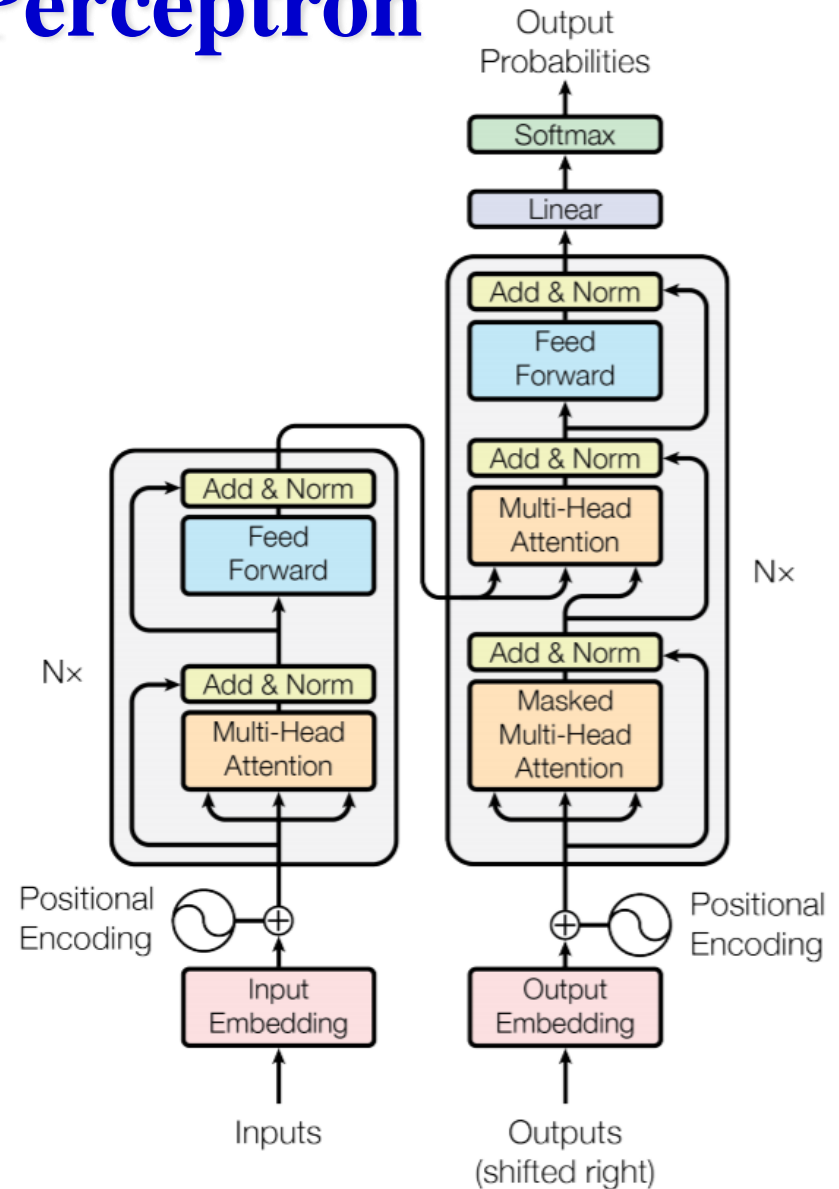
Given the complexity of ML models (i.e., ML functions), it is now common to use a graphical representation of these models that uses higher-level components.



# Multi-Node Perceptron

Here, for example, is a graphical representation of **transformer**, the model that powers ChatGPT.


There are 2 x 96 of these “layers”!



# LOSS FUNCTIONS

# 2D Dataset

- In this part of the course, we'll use a simple synthetic dataset comprising two classes of objects characterized by real numbers ( $x_1, x_2$ ).
- The data are generated from two bivariate normal distributions.




# Loss Functions

Suppose our goal is to approximate an unknown function

$$y = f(x_1, x_2)$$

given a dataset  $D = \{(x_1, x_2, y)_{i=1}^N\}$ ,

perhaps using the *model* on the right.



# Loss Functions

A well-known approach is to approximate  $f(x_1, x_2)$  with a parameterized function  $f(x_1, x_2, \omega)$  and find the *best-fit values*  $\omega^*$  by minimizing

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N (y_i - f_i)^2$$

with respect to the parameters  $\omega$ , where  $f_i = f(x_{1i}, x_{2i}, \omega)$ .

This is often referred to as a **least-squared fit**.

# Loss Functions

In machine learning, one minimizes

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N L(y_i, f_i)$$

where the function,  $L(y_i, f_i)$ , called the **loss function**, measures the discrepancy between the desired **target**  $y$  and the model  $f$ .

The quantity  $R(\omega)$  is called the **empirical risk** and the task of minimizing it is called **empirical risk minimization**.

This is just a generalization of least-squared fitting.



# Risk Functional

In mathematics one can often gain insight by taking the limit of an expression. Let's take the limit of

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N L(y_i, f_i)$$

as  $N \rightarrow \infty$ , that is, as the amount of data grows without limit.

In that limit, the empirical risk becomes the **risk functional**,


$$R[f] = \int dx \int dy L(y, f) p(x, y)$$

where  $p(x, y)dx dy$  is the probability distribution of the data.

# Risk Functional: Landscape

The empirical risk defines a highly corrugated very high-dimensional “landscape” in the space of parameters.

The goal of an **optimizer** is to navigate the landscape ( $R(\omega)$ ) associated with a *finite* amount of data to find a *good approximation* of the lowest point of the landscape ( $R[f]$ ) associated with an *infinite* amount of data. When ML researchers say that a model *generalizes* this is what they mean.



# Summary

- We described the multi-node perceptron and how machine learning models are represented graphically.
- We illustrated the calculation of the output of a perceptron.
- We discussed the generalization of least-squared fitting to empirical risk minimization.