

MACHINE LEARNING IN PHYSICS FOUNDATIONS 2

Harrison B. Prosper
PHY 6938 Fall 2024

Recap: AI, ML, and DL

Artificial Intelligence

Algorithms that cause machines to exhibit human- or *super-human*-level intelligence.

Machine Learning

Algorithms for modeling data


Deep Learning

ML using (large) neural networks

Recap: Deep Learning

Uses functions formed by *composition*:

$$\begin{aligned} f(x) &= f_m \circ f_{m-1} \circ \cdots f_1 \\ &= f_m(f_{m-1}(\cdots f_1(x)) \cdots) \end{aligned}$$



Example:

$$f(x) = \text{softmax} \left(\text{dropout}(\text{linear}(\text{flatten}(\mathbf{g}(\mathbf{c}(\mathbf{h}(\mathbf{c}(x))))))) \right)$$

Recap: The Perceptron


$$y = g(xA^T + b)$$

x is matrix (one or more rows) of input data

A is a matrix of **weights**

b is a matrix of **biases**

$g(*)$ is a nonlinear function



Activations: $g(z) = \text{relu}(z) = \max(0, z)$

$g(z) = \tanh(z)$

$g(z) = \text{sigmoid}(z) = 1 / (1 + \exp(-z))$


<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon>

MULTI-NODE PERCEPTRON

Multi-Node Perceptron

$$g(z) = \text{relu}(z)$$


$$y_1 = g(xA_1^T + b_1)$$



Example:

$$A_1 = (2, -3, 1), \quad b_1 = -5$$


$$y_2 = g(xA_2^T + b_2)$$



$$A_2 = (1, 2, 3), \quad b_2 = -4$$

Multi-Node Perceptron

A multi-node perceptron is often drawn as follows,



computed sequentially:

$$z = xA^T + b$$

$y = g(z)$ and, usually, the

activation function $g(z)$ is applied *elementwise*, that is, to every element of its matrix input. (Demo)

Multi-Node Perceptron


Consider the task of modeling the data triplets $D = \{(x_1, x_2, y)_{i=1}^N\}$ with a function of the form $y = f(x_1, x_2, \omega)$. Let's try a 2-input 4-node perceptron with $g(z) = z\sigma(z)$,

where,

$$\sigma(z) = \frac{1}{(1 + e^{-z})}$$


$z\sigma(z)$ is called a **silu**.

ω are free parameters.




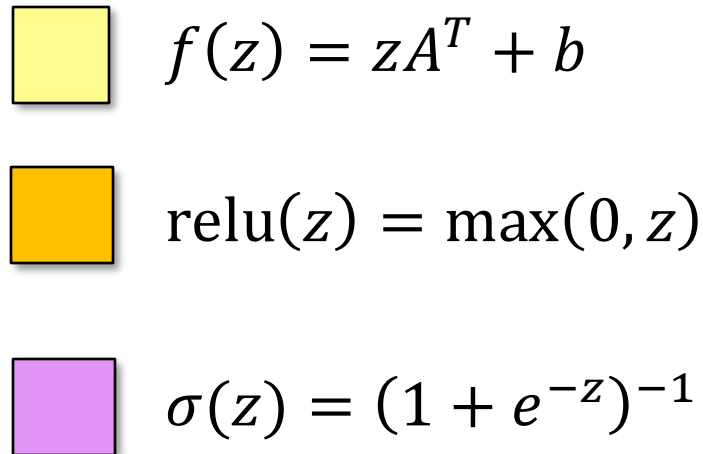
Multi-Node Perceptron

Then, let's feed the output of the first perceptron into a 4-input 1-node perceptron:



Multi-Node Perceptron

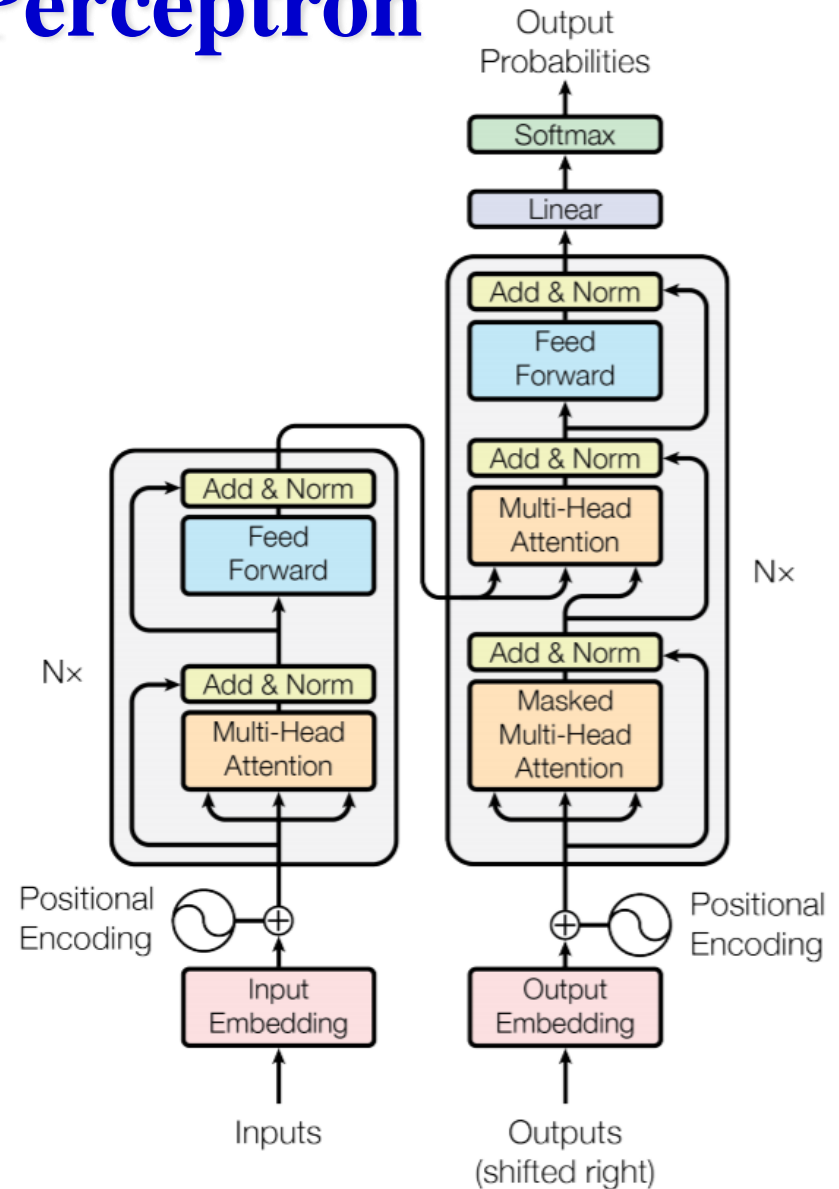
Given the complexity of ML models (i.e., ML functions), it is now common to use a graphical representation of these models that uses higher-level components.



Multi-Node Perceptron

Here, for example, is a graphical representation of a **transformer**, the model that powers ChatGPT.


There are 2 x 96 of these “layers”!



LOSS FUNCTIONS

2D Dataset

- In this part of the course, we'll use a simple synthetic dataset comprising two classes of objects characterized by real numbers (x_1, x_2).
- The data are generated from two bivariate normal distributions.




Loss Functions

Suppose our goal is to approximate an unknown function

$$y = f(x_1, x_2)$$

given a dataset $D = \{(x_1, x_2, y)_{i=1}^N\}$,

using the *model* on the right.



Loss Functions

A well-known approach is to approximate $f(x_1, x_2)$ with a parameterized function $f(x_1, x_2, \omega)$ and find the *best-fit values* ω^* by minimizing

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N (y_i - f_i)^2$$

with respect to the parameters ω , where $f_i = f(x_{1i}, x_{2i}, \omega)$.

This is often referred to as a **least-squared fit**.

Loss Functions

In machine learning, the least-squares fit is generalized to

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N L(y_i, f_i)$$

where the function, $L(y_i, f_i)$, called the **loss function**, measures the discrepancy between the desired **target y** and the model f .

The quantity $R(\omega)$ is called the **empirical risk** and the task of minimizing it is called **empirical risk minimization**.

Warning: In ML, $R(\omega)$ is sometimes referred to as the “loss”.

Risk Functional

In mathematics one can often gain insight by taking the limit of an expression. Let's take the limit of

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N L(y_i, f_i)$$

as $N \rightarrow \infty$, that is, as the amount of data grows without limit.

In that limit, the empirical risk becomes the **risk functional**,


$$R[f] = \int dx \int dy L(y, f) p(x, y)$$

where $p(x, y)dx dy$ is the probability distribution of the data.

Risk Functional: Landscape

The empirical risk defines a highly corrugated very high-dimensional “landscape” in the space of parameters.

The goal of an **optimizer** is to navigate the landscape ($R(\omega)$) associated with a *finite* amount of data to find a *good approximation* of the lowest point of the landscape ($R[f]$) associated with an *infinite* amount of data. When ML researchers say that a model *generalizes* this is what they mean.



Summary

- We described the multi-node perceptron and how machine learning models are represented graphically.
- We illustrated the calculation of the output of a perceptron.
- We discussed the generalization of least-squared fitting to empirical risk minimization.