

MACHINE LEARNING IN PHYSICS FOUNDATIONS 3

HARRISON B. PROSPER

PHY6937 / PHY4936

Recap

What have we learned so far?

1. Machine learning models are trained by minimizing the empirical risk (aka average loss):

$$R(\omega) = \frac{1}{N} \sum_{i=1}^N L(y_i, f_i)$$

2. The loss function, $L(y, f)$, is chosen according to the task at hand.

Recap

3. The empirical risk (which is typically referred to as the “loss” in ML circles) is an *unbiased* Monte Carlo estimate of the risk functional

$$R[f] = \int dx \int dy L(y, f) p(x, y)$$

where $p(x, y)$ is the probability density of the data.

4. Minimizing the risk functional with respect to the ML model f shows that the optimal model satisfies

$$\int dy \frac{\partial L}{\partial f} p(y | x) = 0$$

EMPIRICAL RISK MINIMIZATION


Empirical Risk Minimization (1)

The minimization of $R(\omega)$ is typically done by moving in the direction of *steepest descent*. Most optimization algorithms in ML are variations of **Stochastic Gradient Descent**.

Algorithm

1. At the current point ω_j , compute a *noisy* approximation of the gradient of $R(\omega) \approx \frac{1}{n} \sum_{i=1}^n L(y_i, f_i)$ by using a **batch** of **training data**, where $n \ll N$.
2. Move to the next position ω_{j+1} in the landscape using

$$\omega_{j+1} = \omega_j - \eta \nabla R$$



Empirical Risk Minimization (2)

Why does this (1st-order) algorithm

$$\omega_{j+1} = \omega_j - \eta \nabla R$$


work?

Consider

$$\begin{aligned} R(\omega_{j+1}) &= R(\omega_j - \eta \nabla R) \\ &= R(\omega_j) - \eta \nabla R \cdot \nabla R + O(\eta^2) \end{aligned}$$

If the $O(\eta^2)$ can be neglected, and given that the $O(\eta)$ term is always negative, it follows that

$$R(\omega_{j+1}) < R(\omega_j).$$



BINARY CLASSIFIERS

Binary Classifiers (1)

Starting with

$$\int dy \frac{\partial L}{\partial f} p(y|x) = 0$$

it was shown in Lecture 2 and tutorial 2 that a binary classifier $f(x, \omega)$ trained using the cross-entropy loss $L(y, f) = -[y \log f + (1 - y) \log(1 - f)]$ approximates the **class probability**

$$f(x, \omega^*) = p(1|x)$$

that is, the probability that an object characterized by the **features x** belongs to the class labeled $y = 1$.

Binary Classifiers (2)

From Bayes' theorem, we can write the class probability as

$$p(1|x) = \frac{p(x|1) \pi(1)}{p(x)}$$

where $p(x) = p(x|1) \pi(1) + p(x|0) \pi(0)$ and

- $p(x|k)$ is the probability density of x for class $y = k$,
- $\pi(k)$ is the **prior probability** of class $y = k$, and
- $\pi(1) + \pi(0) = 1$.

Binary Classifiers (3)

When the training dataset is **balanced** (as in tutorial 2), that is, $\pi(1) = \pi(0)$, the class probability $p(1|x)$ becomes the **discriminant**,

$$D(x) = \frac{p(x|1)}{p(x|1) + p(x|0)}$$

This is one of the most useful results in scientific machine learning.

Binary Classifiers (4)

It is possible to train a binary classifier using an **unbalanced** dataset: $\pi(1) \neq \pi(0)$.

But if $\pi(1) \ll \pi(0)$ this might be challenging.

Why do you think this is?)

However, as shown on the next slide, $p(1|x)$ can be computed from a model trained with a balanced dataset.

Binary Classifiers (4)

According to Bayes' theorem

$$p(1|x) = \frac{p(x|1)\pi(1)}{p(x|1)\pi(1) + p(x|0)\pi(0)}$$

Dividing the numerator and denominator by $p(x|1) + p(x|0)$ yields

$$p(1|x) = \frac{D(x)\pi(1)}{D(x)\pi(1) + [1 - D(x)]\pi(0)}$$

So, it is not necessary to use the true priors during training.

CLASSIFIER PERFORMANCE

Classifier Performance (1)

Given a classifier, $D(x)$, the indicator function $I[*]$, the targets y , and a threshold t on $D(x)$, the following counts can be computed:

1. False Negatives (FN): $\sum_{y_i=1} I[D(x_i) \leq t]$
2. True Positives (TP): $\sum_{y_i=1} I[D(x_i) > t]$
3. True Negatives (TN): $\sum_{y_i=0} I[D(x_i) \leq t]$
4. False Positives (FP): $\sum_{y_i=0} I[D(x_i) > t]$


$I[z] = 1$ if z is true, 0 otherwise.

Classifier Performance (2)


FN: $\sum_{y_i=1} I[D(x_i) \leq t] \quad t = 0.5$

TP: $\sum_{y_i=1} I[D(x_i) > t]$

$y = 0$



$y = 1$



TN: $\sum_{y_i=0} I[D(x_i) \leq t]$

FP: $\sum_{y_i=0} I[D(x_i) > t]$

Confusion Matrix

$$\text{FN: } \sum_{y_i=1} I[D(x_i) \leq t]$$

$$\text{TP: } \sum_{y_i=1} I[D(x_i) > t]$$

The **confusion matrix** is a simple way to summarize the counts in the four disjoint regions.




A perfectly separable dataset would yield a diagonal matrix.

Confusion Matrix		0	1
True Labels	0	TN 6860	FP 626
	1	FN 1192	TP 6322
		Predicted Labels	

$$\text{TN: } \sum_{y_i=0} I[D(x_i) \leq t]$$

$$\text{FP: } \sum_{y_i=0} I[D(x_i) > t]$$

Confusion Matrix ($t = 0.9$)



Accuracy, Precision, Recall



Accuracy = $(TP + TN) / (TP + TN + FN + FP)$

Precision = $TP / (TP + FP)$ $P(y = 1 | D > t)$

Recall = $TP / (FN + TP)$

True Positive Rate (TPR) = $TP / (FN + TP)$ $P(D > t | 1)$

False Positive Rate (FPR) = $FP / (TN + FP)$ $P(D > t | 0)$




ROC and AUC

The ROC curve is a plot of
TPR vs. FPR, that is,

$P(D > t|1)$ vs. $P(D > t|0)$

where (usually)

$$D(x) = \frac{p(x|1)}{p(x|1) + p(x|0)}$$



BEYOND CLASSIFICATION

Beyond Classification (1)

Notice that the discriminant

$$D(x) = \frac{p(x|1)}{p(x|1) + p(x|0)}$$

can be rewritten as

$$p(x|1) = \left[\frac{D(x)}{1 - D(x)} \right] p(x|0)$$

which is often referred to as the likelihood ratio trick.

If $p(x|0)$ is a *known* density from which the $y = 0$ vectors are sampled, then the discriminant can be repurposed as a **probability density estimator**.

Beyond Classification (2)

Consider a dataset $U[x]$, of size N , labeled $y = 1$. The goal is to approximate its associated probability density $u(x)$.

Algorithm

1. Generate another dataset $K[x]$, of size N , labeled $y = 0$, where $x \sim k(x) \equiv p(x|0)$, a *known* density*.
2. Train a classifier $D(x)$ using a randomized mixture of the datasets $U[x]$ and $K[x]$.
3. Approximate the unknown density $u(x)$ using

$$u(x) = k(x) \left[\frac{D(x)}{1 - D(x)} \right]$$

* The symbol “ $\sim$ ” means “sampled from”.

Beyond Classification (3)

But we can go further*.

Suppose our dataset U is from a simulation that depends on a set of parameters θ .

Let's sample these parameters from a *known* prior, $\pi(\theta)$. For each sampled point θ , x is sampled from the *unknown* density $u(x|\theta)$ using the simulator.

The dataset, labeled $y = 1$, comprises N pairs $\{(x, \theta)\}$ that constitute a **point cloud** representation of the joint density $u(x, \theta)$.

* Baldi, P., Cranmer, K., Faucett, T. *et al.*
Parameterized neural networks for high-energy physics.
Eur. Phys. J. C **76**, 235 (2016).

Beyond Classification (4)

Algorithm

1. Generate a dataset K of size N , labeled $y = 0$, where θ is sampled from the same prior $\pi(\theta)$. For each θ , sample $x \sim k(x|\theta)$, where $k(x|\theta)$ is a known conditional density.
2. Train a classifier $f(x, \theta, \omega)$ with an admixture of U and K .
3. Compute $u(x, \theta) = k(x|\theta)\pi(\theta) \left[\frac{D(x, \theta)}{1 - D(x, \theta)} \right]$.
4. Divide by $\pi(\theta)$:

$$u(x|\theta) = k(x|\theta) \left[\frac{D(x, \theta)}{1 - D(x, \theta)} \right]$$

Summary

- In addition to the ROC curve and the AUC, the performance of a classifier can be characterized with several other numbers including:
 - False Negatives (FN)
 - True Positives (TP)
 - True Negatives (TN)
 - False Positives (FP)
- A classifier can be repurposed in many ways, including as a probability density estimator.