

MACHINE LEARNING IN PHYSICS

GRAPH NEURAL NETWORKS

HARRISON B. PROSPER

PHY6938

Recap: Normalizing Flows

Let $y = f(x)$ be *vector*-valued with $x, y \in \mathbb{R}^D$. The change of variables formula is

$$p(x) = q(y) |\det J|$$

or

$$\log p(x) = \log q(f(x)) + \log |\det J|$$

where $J = \nabla_x f(x)$ is the Jacobian matrix of the transformation.

A **normalizing flow** is a sequence of **bijections** modeled as neural networks that approximates $y = f(x)$, and its inverse $x = f^{-1}(y)$.

Recap: Diffusion Model

A diffusion model maps a vector x_1 to x_0 . When x_1 is sampled from a D -dimensional Gaussian this induces a sampling of x_0 from the desired distribution, $p(x_0)$. One class of diffusion model entails solving the integro-differential equation,

$$\frac{dx_t}{dt} = \lambda_t x_t + \mu_t q(t, x_t)$$

from $t = 1$ to $t = 0$, where $\lambda_t \equiv \lambda(t)$ and $\mu_t \equiv \mu(t)$ are known functions of “time” and

$$q(t, x_t) = \int_{\mathbb{R}^D} x_0 \frac{p(x_t | x_0)}{p(x_t)} p(x_0) dx_0$$


where $p(x_t | x_0)$ is a D -dimensional Gaussian.

Introduction

1. Convolutional neural networks (CNN)
2. Autoencoder (AE)
3. Physics-informed neural networks (PINN)
4. Flow and diffusion models
5. Graph neural networks (GNN)
6. Transformer neural networks (TNN)


Introduction

Consider the following neutrino event from IceCube.



Introduction

- The data for a single neutrino event can be modeled as a **point cloud**, that is, a cloud of n points in a d -dimensional space.
- Another way to represent a point cloud is as a **graph**.



- A graph $G = (V, E)$ comprises vertices with attributes $V = \{v_i\}$ and edges with attributes $E = \{e_k\}$.


Introduction

- Standard neural networks expect the input data to be of fixed size.
- However, the number of points in a point cloud or the number of vertices in a graph can change from one data instance to the next.
- Using the work of the IceCube collaboration, we'll show how a graph neural network can handle graphs having a varying number of vertices.

GRAPH NEURAL NETWORKS

Graph Neural Networks (GNN)

A graph neural network is designed to operate on data organized as a graph:



A graph enters a **graph processor** on the left and exits the processor on the right with altered graph attributes.

https://github.com/deepmind/graph_nets

IceCube: Data

The data collected by the IceCube detector consists of signals from photomultipliers*.

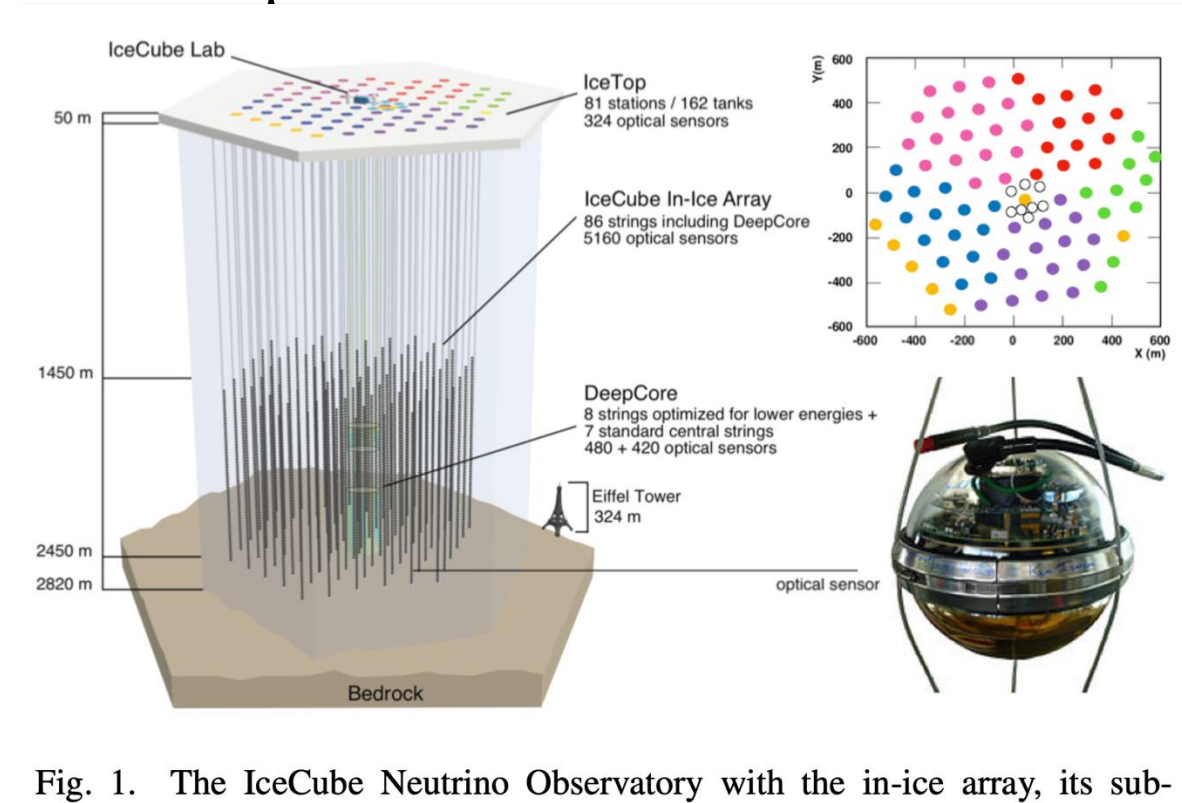



Fig. 1. The IceCube Neutrino Observatory with the in-ice array, its sub-

*N. Choma et al. IceCube collaboration, Graph Neural Networks for IceCube Signal Classification, arXiv:1809.06166v1

IceCube: Data

- IceCube data are the signals from n Digital Optical Modules (DOMs), which are modeled as the vertices of a graph.
- Each vertex is associated with a d -dimensional (row-wise) vector of attributes $v_i = (x_1, \dots, x_d)_i$, three of which are the spatial coordinates (x, y, z) of the DOM.



IceCube: Data


For an event with n signals, an $n \times n$ adjacency matrix, \mathbf{A} , is constructed, as follows, using the spatial data only,

$$(\mathbf{A})_{ij} = \text{softmax}(d_{ij}, \text{dim} = 1),$$

where

$$d_{ij} = \exp\left(-\|x_i - x_j\|^2 / 2\sigma^2\right)$$

and σ is a free parameter. The matrix \mathbf{A} models the edges of the graph.




IceCube: Data

The vectors \mathbf{v}_i are concatenated vertically into an $n \times d$ matrix: $\mathbf{X} = [\mathbf{v}_1; \cdots; \mathbf{v}_n]$

$$\mathbf{X} = \begin{pmatrix} x_1 & y_1 & z_1 & \cdots \\ x_2 & y_2 & z_2 & \cdots \\ \vdots & \vdots & & \vdots \\ x_n & y_n & z_n & \cdots \end{pmatrix}$$

(Note: horizontal concatenation is denoted by $\mathbf{X} = [\mathbf{v}_1, \cdots, \mathbf{v}_n]$.)




IceCube: GNN

- The $n \times d$ matrix, \mathbf{X} , passes through a sequence of T identical graph processors, each with its own parameters.
- A graph processor is parameterized by a $2d \times d/2$ matrix of parameters, \mathbf{w} , and a single scalar b and computes:

$$\mathbf{Y} = [\mathbf{AX}, \mathbf{X}] \mathbf{w} + b \mathbf{I}$$

- The \mathbf{Y} operation is called a **graph convolution (GConv)** and \mathbf{I} is an $n \times d/2$ matrix of ones.




IceCube: GNN

- The vertices, \mathbf{X} , of the output graph $G = (\mathbf{X}, \mathbf{A})$ are constructed through *horizontal* concatenation:

$$\mathbf{X} = [\text{ReLU}(\mathbf{Y}), \mathbf{Y}]$$

with, as usual, the ReLU function applied *element-wise*.



IceCube: GConv

$$Y = [AX, X]w + bI,$$
$$X = [\text{ReLU}(Y), Y]$$

Implementation


$$AX = \text{torch.matmul}(A, X)$$

$$AXX = \text{torch.cat}([AX, X], \text{dim}=1)$$

$$AXXw = \text{torch.matmul}(AXX, w)$$


$$Y = AXXw + b * \text{torch.ones_like}(AXXw)$$

$$X = \text{torch.cat}([\text{F.relu}(Y), Y], \text{dim}=1)$$



IceCube: GNN

- At the end, \mathbf{X} is mapped to a d -dimensional vector \mathbf{x} using a *permutation invariant map* with respect to the vertices and *invariant* with respect to the number of vertices.
- Finally, the vector \mathbf{x} is mapped to the scalar output $0 \leq y \leq 1$ using a *sigmoid*.



IceCube: GNN Results

The GNN is **6.3** times more efficient than the IceCube physics baseline analysis at a signal to noise ratio that is **3** times better.

It also outperforms a 3D CNN.

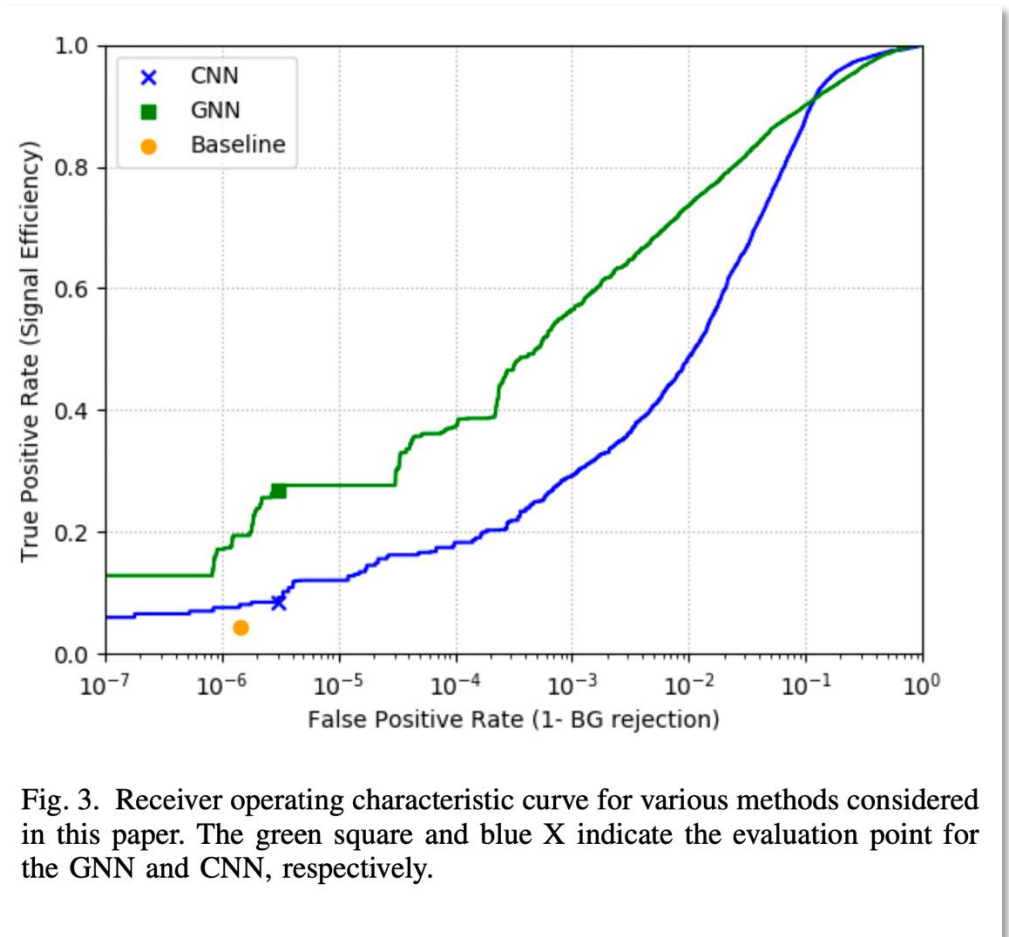


Fig. 3. Receiver operating characteristic curve for various methods considered in this paper. The green square and blue X indicate the evaluation point for the GNN and CNN, respectively.

Summary

- The models we have discussed so far are ill-suited to handle variable input data.
- But data instances in which the number of points, or vertices, can vary can be modeled with graphs $G = (V, E)$ using graph neural networks (GNN).
- We looked at a specific application of GNNs by IceCube, which, for this experiment, achieved state-of-the-art classification results.